PRAJAKTA SHEVAKARI

+1 (443) 401-8792 | prajaktashevakari@gmail.com | Baltimore, USA | Open to relocation | https://www.linkedin.com/in/prajakta-shevakari/

PROFESSIONAL SUMMARY

- Software Engineer with 6+ years of experience designing and building scalable backend and distributed systems to solve complex data and infrastructure challenges in regulated fintech environments.
- Specializes in cloud architecture, microservice design, high-throughput data pipelines, and system reliability at scale.

TECHNICAL SKILLS

- **Backend & Distributed Systems:** Java, Python, Spring Boot, Hibernate, Microservices, REST APIs, Apache Kafka, System Design
- **Cloud & Infrastructure:** AWS (EC2, S3, Lambda, ECS/EKS, ALB/NLB, IAM, CloudWatch)
- **Data Engineering:** Oracle, PL/SQL, SQL, Apache Spark, ETL pipelines, Splunk, Power BI
- **Development & DevOps:** Git, Jenkins, TeamCity, Maven, Chef, CI/CD pipelines, JUnit, Bash, Agile/Scrum
- **Machine Learning (Applied/Integration):** PyTorch, LLMs, Transformers, NLP, ML model integration, LLM inference pipelines

WORK EXPERIENCE

Software Engineer Intern — **June 2025 - August 2025**

Johns Hopkins University - Office of Digital Transformation |Baltimore, MD

- Built Python pipelines serving 200+ faculty members across JHU departments, reducing manual data entry workload by 25%.
- Designed and deployed applications to automate HR data cleanup, streamlining workflows and ensuring 100% data integrity.

Lead Software Engineer — **April 2023 - May 2024**

Barclays | Pune, India

- Architected and built a high-volume executive analytics platform, defining RESTful service contracts and data flows to visualize 1M+ daily transactions in near real-time.
- Led a team of 2 engineers, conducting code reviews, mentoring junior developers, and coordinating with product, QA, and infrastructure teams to define technical standards and delivery processes.
- Automated DR deployments via Autosys, TeamCity, and Chef, cutting release latency from 4+ hours to under 10 minutes.

Software Engineer III — **October 2021 - April 2023**

JPMorgan Chase | Mumbai, India

- Led cloud migration of terabyte-scale regulated data to AWS (S3, EC2), driving architectural decisions across engineering, compliance, and infrastructure teams. Built a high-throughput enrichment engine using Java and Apache Spark, ingesting GBs daily for compliance reporting with zero data loss.
- Redesigned legacy monoliths into resilient microservices, reducing MTTR by 70%. Defined service boundaries and data contracts with principal architects for the firm's hybrid-cloud platform.

Software Engineer — **July 2017 - October 2021**

FIS Global | Pune, India

- Managed SDLC for 90+ weekly production deployments using Jenkins CI/CD pipelines, achieving a 95% on-time delivery rate by prioritizing showstopper defects.
- Built a bulk data upload tool using Spring Boot and Oracle to automate reference data ingestion, eliminating 10+ hours/week of manual operations and reducing data entry errors by ~90%; recognized with an Above and Beyond award.

EDUCATION

Johns Hopkins University — **August 2024 - May 2026**

Master's, Computer Science — **GPA: 3.7**

- Artificial Intelligence, Machine Learning, Natural Language Processing, ML: AI System Design & Development, Future of Work with AI, Human Language Technology, Software System Design, Parallel Computing for Data Science

Savitribai Phule Pune University — **August 2013 - May 2017**

Bachelor's, Computer Engineering — **GPA: 3.8**

PROJECTS

D-FuseSmileNet: Spontaneous Smile Recognition | PyTorch, Transformers — 2024

- Built PyTorch inference pipeline with parallel feature extraction, benchmarking 15 fusion strategies across four datasets.
- Designed reusable data ingestion and evaluation components for consistent cross-dataset training and testing on AWS EC2.

Clinical NER with LLMs and Chain-of-Thought (CoT) Prompting | GPT-4-mini, Hugging Face — 2025

- Built automated annotation pipeline using LLM prompting (few-shot, chain-of-thought) to generate synthetic training data from 144K+ clinical sentences.
- Integrated preprocessing, fine-tuning, and evaluation into a single workflow, replacing manual annotation across three transformer-based models.

TerraNova: Post-Wildfire Decision Support | U-Net, Random Forest, Genetic Algorithm — 2025

- Built backend pipelines integrating satellite imagery, geospatial data, and ML models (U-Net segmentation, reburn-risk prediction) for post-wildfire assessment.
- Developed data services and APIs to support scenario-based analysis through an interactive web interface.